

Querying Data on Decentralized Networks

Matteo Campanelli

Protocol Labs

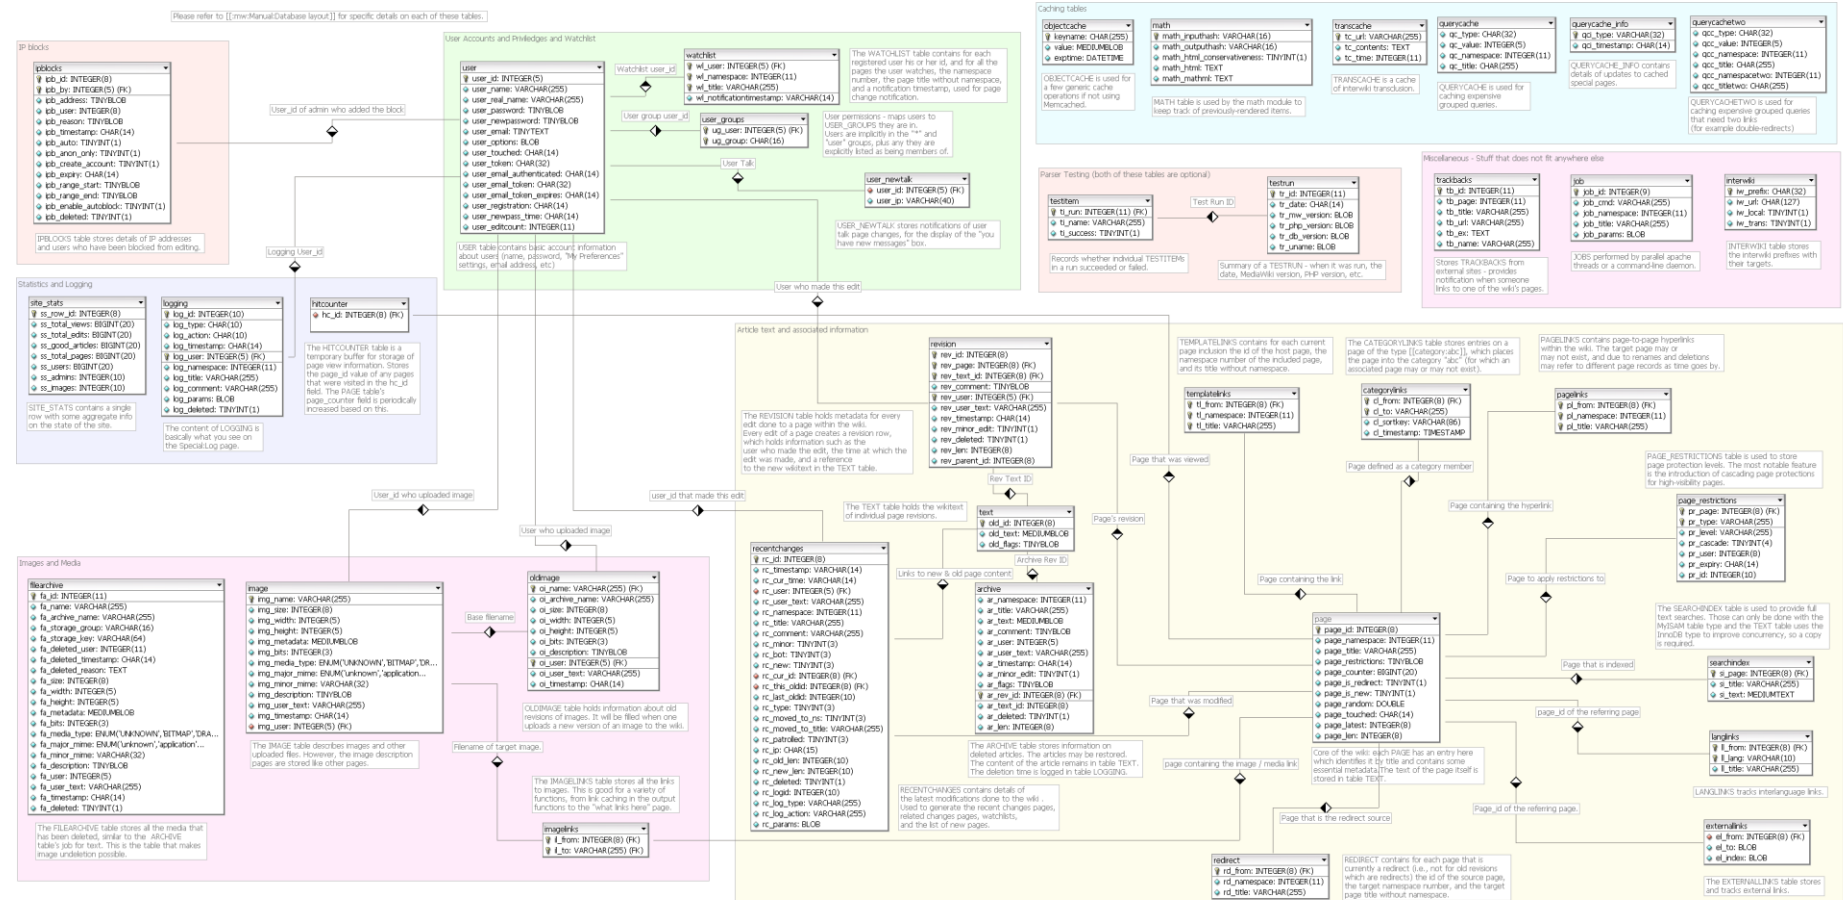
Meta: A link to these slides

- At: binarywhales.com/EC22.pdf
 - They contain additional pointers, notes and references

This talk at the high level:

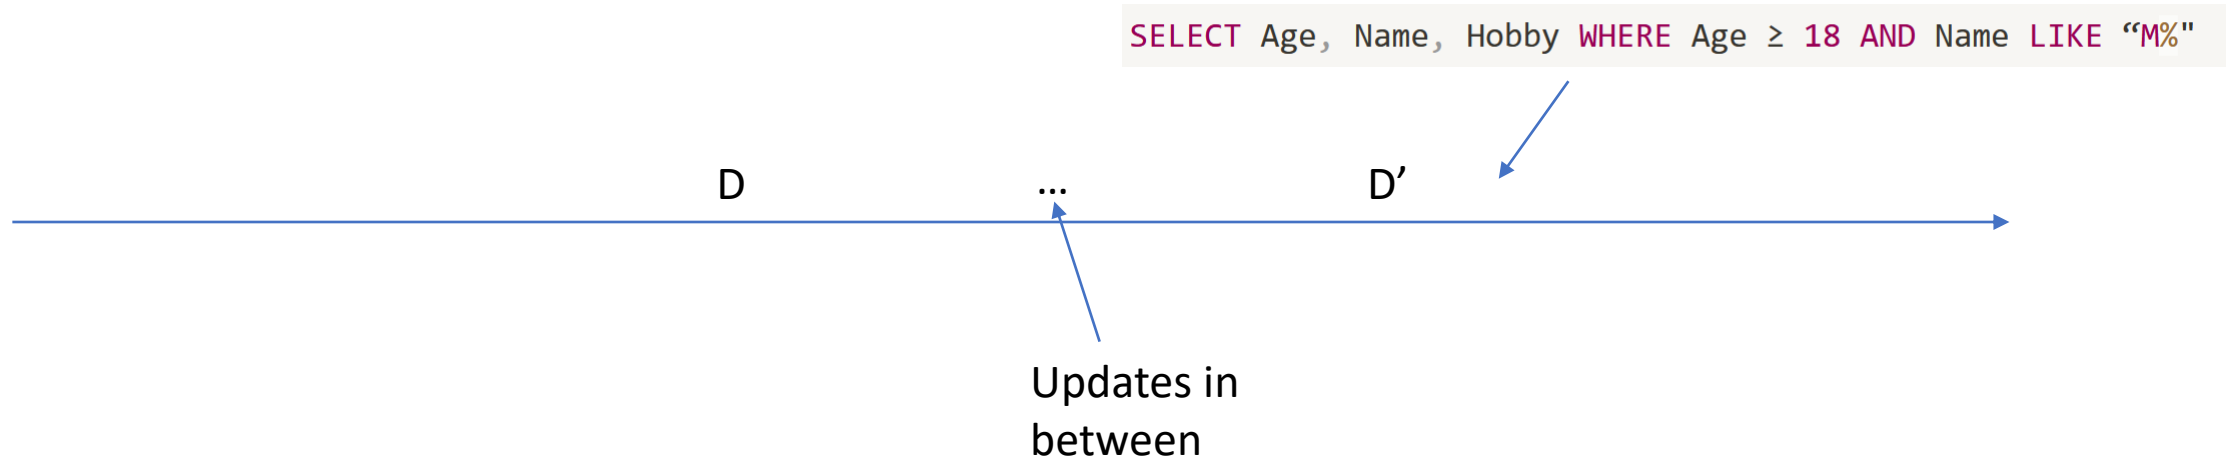
- Here is an application
- Unclear if what is out there is the best (or good) solution for it

Data are at the core of applications



A schema of Wikipedia DB. Source: Wikipedia

Data in applications traditionally



Data and queries in this talk

- Queries on data that are:
 - On chain*
 - Verifiable

*on chain ~ on decentralized networks

On chain data




$[x]$ = “small digest to x ”

Change 1:

- Chain storage is expensive
- no actual data on chain;
just digest

On chain data



$[x]$ = "small digest to x"

Change 2:

- We cannot trust who delivers the query response
- Add **Publicly Verifiable** proof of correct query execution
 - $\text{Verify}([D'], \text{Query}, \text{Result}, \text{Proof}) \rightarrow \text{accept/reject}$

Why?

- Smart contracts can verifiably query DBs
 - E.g. “Update internal state with (verified) result of query Q”
 - E.g. “Release coins to the result of query*”
`SELECT ARGMAX(effectiveness_score) FROM organizations WHERE type=charity`
- In general: achieve complex logic & state with succinctness and verifiability

*Not real SQL syntax, but you get the gist.

Question of this talk: *How to build this?*

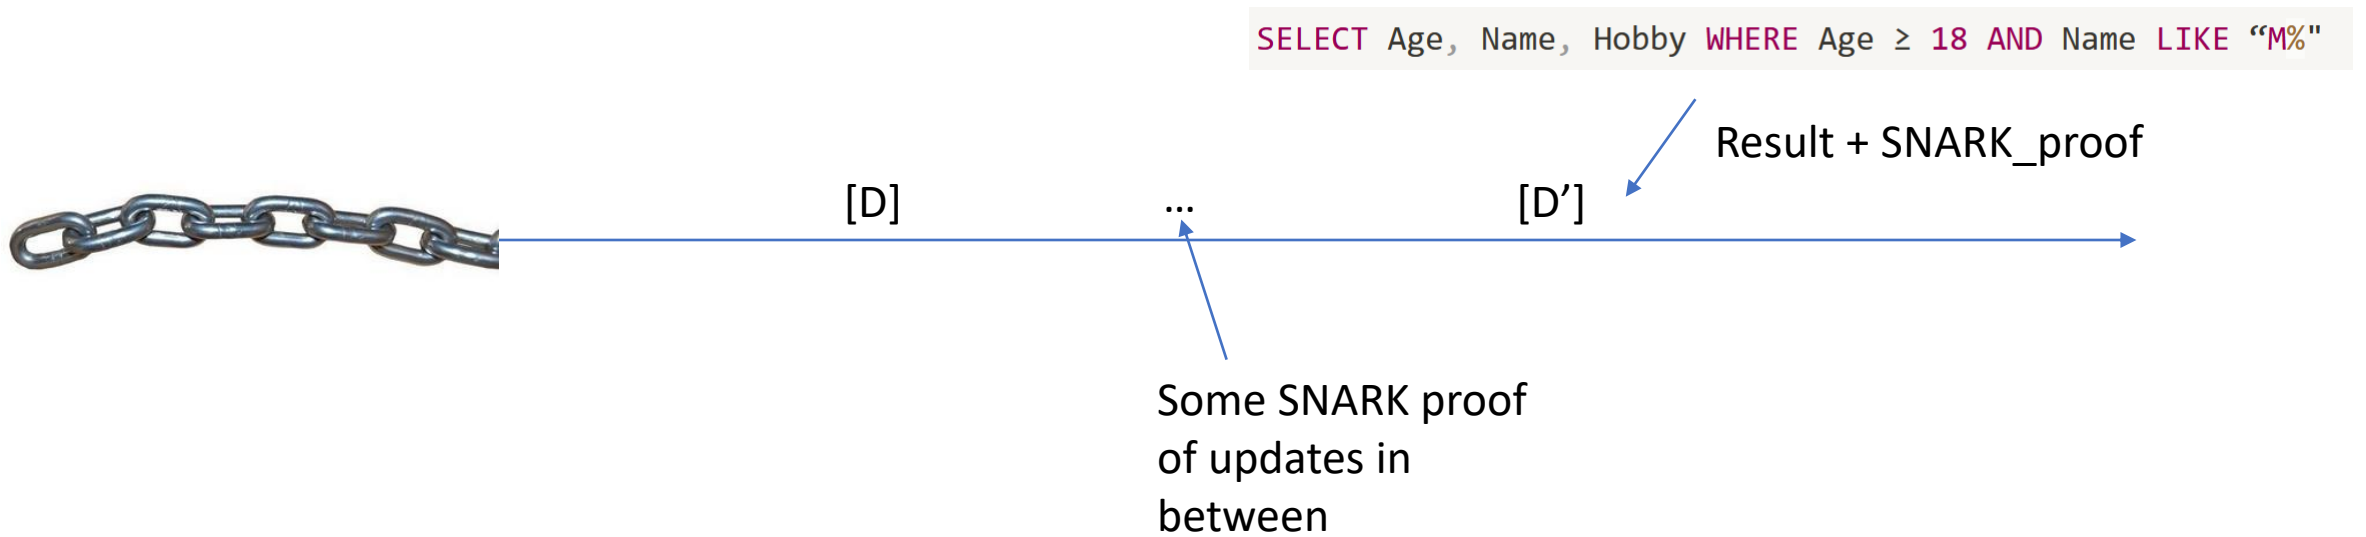
- There are cryptographic primitives that are available
- But it is unclear* if any of them is right for the job

* At least, it is unclear to yours truly

Let's start:

what's the first candidate we could think of?

SNARKs



[x] = "small digest to x"

SNARKs as candidate for the job

Some features of SNARKs

+ expressive (any query)

+ good time(V) and $|\text{prf}|$

- **Too general?**

- Prover has high cost

The (potential) issue with “too general”

- Conversions of computations into circuits is generally expensive
- Conversions of computations into circuits/RAM/etc is generally messy
- **There are many lost opportunities for optimizations**

Going around “too general”

Option #1: engineering SNARKs appropriately

- **Research Question A:**

What are ways we can optimize general SNARKs for a database-like setting?

vSQL: Verifying Arbitrary SQL Queries over
Dynamic Outsourced Databases

Yupeng Zhang*, Daniel Genkin^{†,*}, Jonathan Katz*, Dimitrios Papadopoulos^{‡,*} and Charalampos Papamanthou*
*University of Maryland [†]University of Pennsylvania [‡]Hong Kong University of Science and Technology
Email: {zhangyp,cpap}@umd.edu, danielg3@cis.upenn.edu, jkatz@cs.umd.edu, dipapado@cse.ust.hk

Nova: Recursive Zero-Knowledge Arguments
from Folding Schemes

Abhiram Kothapalli^{*†} Srinath Setty* Ioanna Tzialla[‡]
*Microsoft Research [†]Carnegie Mellon University [‡]New York University


LegoSNARK: Modular Design and Composition of
Succinct Zero-Knowledge Proofs

Matteo Campanelli¹, Dario Fiore¹, and Anaïs Querol^{1,2}

Going around “too general”

Option #2: let’s think “non-general and non-SNARKs”

- **Warm-up question.**



Succinct Functional Commitments (sFC)

Authenticated Data Structures

- Very active area earlier in the millennium (~ 2002-2013)
- **Goal:** ~ making cryptographic version of data structures
 - Think: interval trees, dictionaries, etc.
- **The question for us:** What can they bring to the table here?
- **Challenges (?)**:
 - "Different" security definition (e.g. rely on trusted updates/generation)
 - Most of them are very specific to certain tasks

Streaming Authenticated Data Structures

Charalampos Papamanthou¹, Elaine Shi², Roberto Tamassia³, and Ke Yi⁴

Efficient Authenticated Data Structures for Graph Connectivity and Geometric Search Problems*

MICHAEL T. GOODRICH[†] ROBERTO TAMASSIA[‡] NIKOS TRIANOPOULOS^{‡,§}

- **Research Question B:**

How can we leverage existing constructions/techniques in Auth Data Structures (and sFC) for our goal?

Subquestions:

Are there limitations in their security models for our setting and how to go around them?
When are they really more efficient?

- **(Meta) Research Question C*:**
What is the minimal set of “non general” queries that would be worth having in applications?

* Support question for question B

That's all! Questions?

- **Research Question A:**
What are ways we can optimize general SNARKs for a database-like setting?
- **Research Question B:**
How can we leverage existing constructions/techniques in Auth Data Structures (and sFC) for our goal?
- **Research Question C:**
What is the minimal set of “non general” queries that would be worth having in applications?

These slides (with additional pointers) available at:
binarywhales.com/EC22.pdf

For questions/comments:
matteo@protocol.ai

Some pointers on Authenticated Data Structures

- IntegriDB: this system is the closest thing to something that could be used in practice. Its problem seems to be the requirement (as in other ADS) of a secret key for certifying the updates. We do not expect to have a secret key in our setting
- The two works it is based on
 - <https://eprint.iacr.org/2010/455.pdf>
 - <https://eprint.iacr.org/2013/724.pdf>
- Streaming Authenticated Data Structures (2013):
<http://elaineshi.com/docs/streaming.pdf>

On existing solutions

Auth Data Structures	Succinct Functional Commitments	Snarky approaches*
+ vast literature	+ “nice(r)” assumptions (than SNARKs)	+ expressive (any query)
+ simple constructions	+ (some) simple constructions	+ extractability
? Succinct updates?	+ succinctness / non det	+ good time(V) and prf
? Security model for non-det ?	? When is eval binding sufficient for application(vs, say, extractability)s ?	- Prover is high cost
- Limited to specific queries	? Efficient in practice ?	- Too generic? Losing opportunities for optimizations?
	? Expressive queries ?	- Less “nice” Assumptions (comparatively)

*includes vSQL and recursive approaches

Why care about better proving time

- Counterarguments to “Just delegate to somebody with more powerful hardware”
 - Privacy
 - Democratizing as much as possible roles in decentralized networks (o.w. this is introducing yet another plutocracy)
 - Even powerful hardware is going to hit a point of max capacity at some point. Let's lower that point with better proving time
 - The “Why not?”-response: lower proving time is a better dimension for proof systems

Security of (streaming) ADS

Definition 3 (Security). Let \mathcal{A} be an SADS scheme consisting of the set of algorithms $\{\text{genkey}, \text{initialize}, \text{updateVerifier}, \text{updateProver}, \text{query}, \text{verify}\}$, k be the security parameter, D_0 be the empty data structure and $\text{pk} \leftarrow \text{genkey}(1^k)$. Let also Adv be a PPT adversary and let d_0 be the state output by $\text{initialize}(D_0, \text{pk})$.

- (Update) For $i = 0, \dots, h - 1 = \text{poly}(k)$, Adv picks the update upd_i to data structure D_i . Let $d_{i+1} \leftarrow \text{updateVerifier}(\text{upd}_i, d_i, \text{pk})$ be the new state corresponding to the updated data structure D_{i+1} .
- (Forge) Adv outputs a query q , an answer α and a proof Π .

We say that the SADS scheme \mathcal{A} is secure if for all $k \in \mathbb{N}$, for all pk output by algorithm genkey , and for any PPT adversary Adv it holds that

$$\Pr \left[\begin{array}{l} \{q, \Pi, \alpha\} \leftarrow \text{Adv}(1^k, \text{pk}); 1 \leftarrow \text{verify}(q, \alpha, \Pi, d_h, \text{pk}); \\ 0 \leftarrow \text{check}(q, \alpha, D_h). \end{array} \right] \leq \text{neg}(k). \quad (2.1)$$

Streaming Authenticated Data Structures

Piperine

Replicated state machines without replicated execution

Jonathan Lee Kirill Nikitin* Srinath Setty
*Microsoft Research *EPFL*