

Towards a simple modern
design for VDB

Some VDB papers

Database Outsourcing with Hierarchical Authenticated Data Structures

Mohammad Etemad Alptekin Küpçü
Crypto Group, Koç University, Istanbul, Turkey
{metemad, akupcu}@ku.edu.tr

Super-Efficient Verification of Dynamic Outsourced Databases*

Michael T. Goodrich¹, Roberto Tamassia², and Nikos Triandopoulos³

¹ Dept. of Computer Science, UC Irvine, USA, goodrich@ics.uci.edu

² Dept. of Computer Science, Brown University, USA, rt@cs.brown.edu

³ Dept. of Computer Science, University of Aarhus, Denmark, nikos@daimi.au.dk

Abstract. We develop new algorithmic and cryptographic techniques for authenticating the results of queries over databases that are outsourced to an untrusted responder. We depart from previous approaches by considering *super-efficient* answer verification, where answers to queries are validated in time asymptotically less than the time spent to produce them and using lightweight cryptographic operations. We achieve this property by adopting the decoupling of query answering and answer verification in a way designed for queries related to range search. Our techniques allow for efficient updates of the database and protect against replay attacks performed by the responder. One such technique uses an

vSQL: Verifying Arbitrary SQL Queries over Dynamic Outsourced Databases

Yupeng Zhang*, Daniel Genkin^{†,*}, Jonathan Katz*, Dimitrios Papadopoulos^{‡,*} and Charalampos Papamanthou*

*University of Maryland [†]University of Pennsylvania [‡]Hong Kong University of Science and Technology

Email: {zhangyp,cpap}@umd.edu, danielg3@cis.upenn.edu, jkatz@cs.umd.edu, dipapado@cse.ust.hk

Abstract—Cloud database systems such as Amazon RDS or Google Cloud SQL enable the outsourcing of a large database to a server who then responds to SQL queries. A natural problem here is to efficiently verify the correctness of responses returned

ever, they suffer from limited expressiveness, and in particular they cannot handle a wide range of SQL queries.

A. Our Results

IntegrIDB: Verifiable SQL for Outsourced Databases

Yupeng Zhang
ECE Dept. & UMIACS
University of Maryland
zhangyp@umd.edu

Jonathan Katz
CS Dept. & UMIACS
University of Maryland
jkatz@cs.umd.edu

Charalampos Papamanthou
ECE Dept. & UMIACS
University of Maryland
cpap@umd.edu

ABSTRACT

This paper presents INTEGRIDB, a system allowing a data owner to outsource storage of a database to an untrusted server, and then enable anyone to perform verifiable SQL queries over that database.

most prevalent applications running in the cloud today, and could be offered as a software layer on top of any SQL implementation.

Toward this end we design, build, and evaluate INTEGRIDB, a system that efficiently supports verifiability of a rich subset of SQL

Design Criteria—What we'd like

- Practical
 - Concrete efficiency
 - Conceptual simplicity
 - Development simplicity
 - (example benchmark: I should be able to program it in a few focused days)
 - The "right" level of expressivity
 - It won't need to be almighty
- Modern
 - should potentially use tricks from the last few years
 - E.g., (potentially) aggregations SnarkPack-style, subvector commitments, SNARKs in few localized places, potentially Caulk, etc.

Example DB

Source of table: IntegriDB paper

| row_ID | student_ID | age | GPA | First_name | Type |
|--------|------------|-----|-----|------------|-------|
| 1 | 10747 | 22 | 3.5 | Bob | M.Sc |
| 2 | 10715 | 24 | 3.3 | Alice | B. Sc |
| 3 | 10721 | 23 | 3.7 | David | Ph.D |
| 4 | 10781 | 21 | 3.0 | Cathy | M.Sc |

Table “Students”

| student_ID | title | topic |
|------------|---------------------------------|--------|
| 10721 | “Tarantino and critical theory” | cinema |
| 10715 | “An intro to homotopy theory” | math |
| 10801 | “Lars Von Trier and Dogma 95” | cinema |

Table “FinalWriteups”

Simple query example

| row_ID | student_ID | age | GPA | First_name |
|--------|------------|-----|-----|------------|
| 1 | 10747 | 22 | 3.5 | Bob |
| 2 | 10715 | 24 | 3.3 | Alice |
| 3 | 10721 | 23 | 3.7 | David |
| 4 | 10781 | 21 | 3.0 | Cathy |

Table Students

```
SELECT  
* FROM Students WHERE  
Age IN (21,22,23)
```

Simple JOIN query

| row_ID | student_ID | age | GPA | First_name | Type |
|--------|------------|-----|-----|------------|-------|
| 1 | 10747 | 22 | 3.5 | Bob | M.Sc |
| 2 | 10715 | 24 | 3.3 | Alice | B. Sc |
| 3 | 10721 | 23 | 3.7 | David | Ph.D |
| 4 | 10781 | 21 | 3.0 | Cathy | M.Sc |

Table “Students”

| student_ID | title | topic |
|------------|---------------------------------|--------|
| 10721 | “Tarantino and critical theory” | cinema |
| 10715 | “An intro to homotopy theory” | math |
| 10801 | “Lars Von Trier and Dogma 95” | cinema |

Table “FinalWriteups” (FW)

```
SELECT
Students.student_ID,
Students.GPA,
B.topic
FROM Students INNER JOIN FW
ON Students.student_ID = FW.student_ID
```

Features of the construction

- + good with categorical data (MSc/BSc/PhD, small ranges)
- “Simple” queries
- Think: static databases for now
 - We can think of extending it later
 - NB: static DB == setting with observed preprocessing
- Constant size openings
- Main ingredients:
 - Vector commitments
 - Accumulators with some special properties
 - We’ll see them later

First Attempt: one VC per column

| row_ID | student_ID | age | GPA | First_name |
|--------|------------|-----|-----|------------|
| 1 | 10747 | 22 | 3.5 | Bob |
| 2 | 10715 | 24 | 3.3 | Alice |
| 3 | 10721 | 23 | 3.7 | David |
| 4 | 10781 | 21 | 3.0 | Cathy |

When do we build this list of VCs?

Think at preprocessing time (again,
for now it's a static DB, or a DB where
all updates are observed and
preprocessed honestly)

VC was not enough; let's add something else

| row_ID | student_ID | age | GPA | First_name |
|--------|------------|-----|-----|------------|
| 1 | 10747 | 22 | 3.5 | Bob |
| 2 | 10715 | 24 | 3.3 | Alice |
| 3 | 10721 | 23 | 3.7 | David |
| 4 | 10781 | 21 | 3.0 | Cathy |

Table Students

```
SELECT  
* FROM Students WHERE  
Age = 21
```

VC was not enough b/c we did not know if there were giving **all** rows satisfying a certain property
(VC can only tell you “this subset is in the col”, not necessarily the *maximal* subset)

Idea: Since we are preprocessing the DB, why not having a preprocessing that helps us a little bit with this?

Let's add an "Oracle"

- Before we look at the Auth. Data Structure corresponding to this, let's assume we have a magic oracle for the following
- Let's then see if oracle + VC give us query responses that make sense

The Oracle $\mathbf{O}(T,C,v)$:

- Given table T, column C, value v
- Return { id: $T.C[id] = v$ } // id dependent on ctx (think: some key)

| student_ID | title | topic |
|------------|---------------------------------|--------|
| 10721 | "Tarantino and critical theory" | cinema |
| 10715 | "An intro to homotopy theory" | math |
| 10801 | "Lars Von Trier and Dogma 95" | cinema |

Table "FinalWriteups" (FW)

Example:

$\mathbf{O}(\text{FW}, \text{topic}, \text{"cinema"}) \rightarrow$
{ 10721, 10801 }

Simple query with oracle + VC

| student_ID | title | topic |
|------------|---------------------------------|--------|
| 10721 | "Tarantino and critical theory" | cinema |
| 10715 | "An intro to homotopy theory" | math |
| 10801 | "Lars Von Trier and Dogma 95" | cinema |

Table "FinalWriteups" (FW)

```
SELECT
* FROM FW
WHERE topic ="cinema"
```

Verifier:

- Receive results **Rslt**
- Asks $\mathbf{O}(\text{FW}, \text{topic}, \text{"cinema"})$ // we implement this later as ADS; no worries
- Receive response **Idxs_oracle**
- Receive VC proofs from prover for each col in **Rslt** (see below)
- For each column \$Col in **Rslt**:
 - VC.Vfy(vc: VC_FW_\$Col,
prf: prf_\$Col,
vals: **Rslt**[\$Col],
Idxs: Idxs_oracle)

Recall:

$\mathbf{O}(\text{FW}, \text{topic}, \text{"cinema"}) \rightarrow$
 $\{ 10721, 10801 \}$

Implementing the oracle

The Oracle $\mathbf{O}(T,C,v)$:

- Given table T, column C, value v
- Return S_v // NB: it's a set
- - where $S_v := \{ \text{id}: T.C[\text{id}] = v \}$

The Data Structure for $\mathbf{O}(T,C,v)$:

- Given table T, column C, have a “VC” such that
- $VC = VC.Commit([v1 : Acc_v1, v2 : Acc_v2, \dots])$ // it's more a KV map
- - where $Acc_v := Accum(\{ \text{id}: T.C[\text{id}] = v \})$

Loose ends: size of parameters. A problem?

- What I presented so far has a “digest” linear in the total number of columns.
 - One VC per column + one “oracle ADS” per column
- Is that a problem?

How to do JOINS

Table “Students”

| row_ID | student_ID | age | GPA | First_name | Type |
|--------|------------|-----|-----|------------|-------|
| 1 | 10747 | 22 | 3.5 | Bob | M.Sc |
| 2 | 10715 | 24 | 3.3 | Alice | B. Sc |
| 3 | 10721 | 23 | 3.7 | David | B.Sc |
| 4 | 10781 | 21 | 3.0 | Cathy | M.Sc |

Table “FinalWriteups” (FW)

| student_ID | title | topic |
|------------|---------------------------------|--------|
| 10721 | “Tarantino and critical theory” | cinema |
| 10715 | “An intro to homotopy theory” | math |
| 10801 | “Lars Von Trier and Dogma 95” | cinema |

Verifier:

- Receive results **Rslt**
- Asks $\mathbf{O}(\text{Students}, \text{type}, \text{B.Sc})$
- Asks $\mathbf{O}(\text{FW}, \text{topic}, *)$ // all topics
- Receive responses `Idxs_oracle_Students`, `Idx_oracle_FW`
- Check that `Rslt[id] == Idxs_oracle_Students \cap Idx_oracle_FW` // new op!
- Receive VC proofs from prover for each table/col in **Rslt** (see below)
- For each table \$T and each column \$Col in **Rslt**:
 - VC.Vfy(vc: VC_{\$T}_{\$Col},
prf: prf_{\$T}_{\$Col},
vals: **Rslt**[\$Col],
Idxs: Idxs_oracle)

```
SELECT  
S.Id, S.Name, topic FROM  
S INNER JOIN FW  
WHERE S.type = “B.Sc”
```

How to implement intersections?

- Check that $\text{Rslt}[\text{id}] == \text{Idxs_oracle_Students} \cap \text{Idx_oracle_FW}$ // new op!
- Recall this "oracle" is implemented as accumulators
- We need accumulators supporting intersection
- KZG-based accumulators have this property!
 - KZG accum = Commit(characteristic polynomial of set)
 - Use divisibility for intersections

Other loose ends

- What I am looking at:
 - Sorting gadgets (useful on their own and for updates)
 - VCs are often homomorphic. Let's exploit that!
 - Gadgets to prove:
SELECT * FROM C1, C2 WHERE $C1 + 2 * C2 = 0$
SELECT * FROM C1, C2 WHERE $C1 + 2 * C2 > 0$

-